

Deep learning for NLP

Eneko Agirre, Gorka Azkune
Ander Barrena, Oier Lopez de Lacalle

@eagirre @gazkune @4nderB @oierldl #dl4nlp

<http://ixa2.si.ehu.eus/eneko/dl4nlp>

Session 2: Multilayer Perceptron



Plan for the course

- Introduction: machine learning and NLP
- Multilayer perceptron
- Word representation and Recurrent neural networks (RNN)
- Sequence-to-Sequence (seq2seq) and Machine Translation
- Attention, transformers and Natural language inference
- Pre-trained transformers, BERT, GPT
- Bridging the gap between natural languages and the visual world



Quiz

Find definition and slide for the following:

- Supervised machine learning
- Document classification
- Document regression
- Linear regression
- Logistic regression
- Train, development, test
- Softmax classification
- Loss function J
- Gradients ∇
- Stochastic gradient descent
- Learning rate η
- Mini-batch
- Optimizer
- Overfitting
- Regularization, L2, early stopping




Plan for this session

- Multiple layers ~ Deep: MLP
- Learning rate
- More regularization
- Hyperparameters
- Backpropagation and gradients


Deep: Multilayer perceptron

Logistic Regression



- An input layer – just a feature vector
- An output layer – class probabilities

Deep: Multilayer perceptron




- An input layer
 - just a feature vector
- One or more hidden layers, each computed on the layer below
 - latent features.
- An output layer, based on the top hidden layer
 - class probabilities
- Also known as **Feed Forward** or **Dense**


Deep: Multilayer perceptron

Logistic Regression

$$y = \text{softmax}(xW + b)$$



Deep: Multilayer perceptron




$$y = \text{softmax}(h_1 W_2 + b_2)$$

$$h_1 = f(h_0 W_1 + b_1)$$

$$h_0 = f(x W_0 + b_0)$$

Deep: Multilayer perceptron



$$y = \text{softmax}(h_1 W_2 + b_2)$$

$$h_1 = f(h_0 W_1 + b_1)$$

$$h_0 = f(x W_0 + b_0)$$

softmax

$$J_x = -\log P(y=c|x) = -\log \left(\frac{\exp(h_1 W_2[c])}{\sum_{c' \in C} \exp(h_1 W_2[c'])} \right)$$

Deep: Multilayer perceptron

Layers have the same structure


$$h_i = f(h_{i-1} W_i + b_i)$$

Non-linear functions!

$$\text{Sigmoid: } \sigma(x) = \frac{1}{1 + \exp(-x)}$$




$$\text{Hyperbolic: } \tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$$

$$\text{Rectified linear unit: } \text{rect}(x) = \max(0, x)$$



Deep: Multilayer perceptron


Motivations?



Source: <http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>

Without non-linearities, there is no extra expresivity: a sequence of linear transformations is a linear transformation

Deep: Multilayer perceptron




A MLP with one layer can learn to reproduce any function

- Non-linear functions!

Why should we need anything else?

Deep: Multilayer perceptron

Training: SGD (again!)



- Start with random parameters: W (includes bias)
- Each epoch
 - Shuffle training data
 - For each mini-batch (set of K examples)
 - **Compute the loss function (forward)**
 - **Compute the gradient of the loss function (backward)**
 - Update parameters:
(learning rate η)
$$W = W - \eta \frac{1}{K} \sum_i^{K-1} \nabla J_i(W)$$
 - Measure train and dev. accuracy
- Continue until loss function converges / time is up / dev. accuracy stops increasing

Supervised doc. Classification

Softmax classification

Overfitting and regularization

- W can be very good for training, with enough layers and capacity the model can memorize the training data!
 - Generalize very poorly to test data (= the real world)
- First solution: add a **regularizer** to the loss function that avoids the model to fit the training data

$$J_i(W) = -\log \left(\frac{\exp(W_{c_i}^T x)}{\sum_{c' \in C} \exp(W_{c'}^T x)} \right) + \lambda \sum_k W_k^2$$

squared L2 norm




Supervised doc. Classification

Softmax classification

Overfitting and regularization

- Overfitting can be seen in this graph
- **Early stopping** finishes training as soon as development error starts to increase
- **Experimental setup:** %80 train, %10 development, %10 test (blind!!)
- **Model selection:** best accuracy (lowest error) at development




Source: chatbotlife.com


Plan for this session

- Multiple layers ~ Deep: MLP
- Learning rate
- More regularization
- Hyperparameters
- Backpropagation and gradients

SGD convergence: learning rate




$$W = W - \eta \frac{1}{K} \sum_i^{K-1} \nabla J_i(W)$$



Source: [cs231n.github.io](https://github.com/cs231n)

SGD convergence: learning rate




SGD convergence

Optimizers use diff. learning rates

- Fixed, adam, ...

Don't expect convergence: use early stopping

Don't expect a unique solution: ensembling helps



$$W = W - \eta \frac{1}{K} \sum_i^{K-1} \nabla J_i(W)$$



Plan for this session

- Multiple layers ~ Deep: MLP
- Learning rate
- More regularization
- Hyperparameters
- Backpropagation and gradients



Overcapacity and regularization

The more layers / hidden units, the more capacity and risk of overfitting (memorize whole training data)!




Source: chatbotslife.com

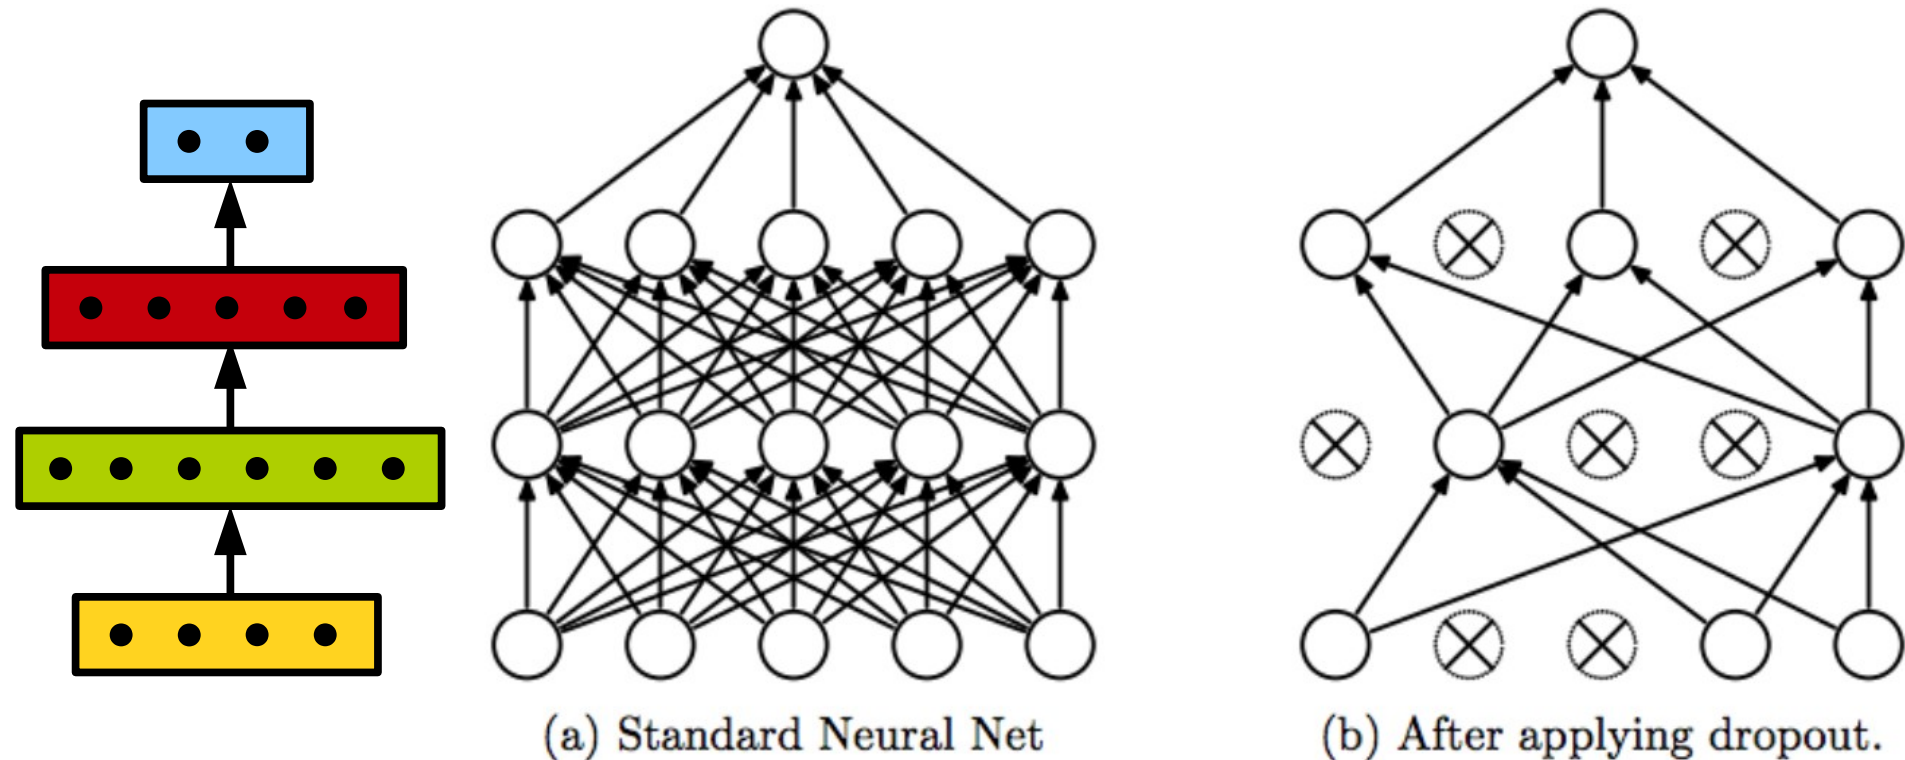
Overcapacity and regularization

The more layers / hidden units, the more capacity and risk of overfitting (memorize whole training data)!

- L2 on parameters $J_i(W) = \dots + \lambda \sum_k W_k^2$
- Early stopping
- Dropout
 - At training time deactivate 50% of the activations at random
 - At test time use all activations




Overcapacity and regularization



Source: “Dropout: a simple way to prevent neural networks from overfitting”, JMLR 2014

MLP: hyperparameters



Topology: number and size of layers

Non-linearity

Optimizer

Learning-rate

Size of mini-batch

Weight of L2 regularization

Dropout rate

Quiz

Find definition and slide for the following:

- Supervised machine learning
- Document classification
- Document regression
- Linear regression
- Logistic regression
- Train, development, test
- Softmax classification
- Loss function J
- Gradients ∇
- Stochastic gradient descent
- Learning rate η
- Mini-batch
- Optimizer
- Overfitting
- Regularization, L2, early stopping



Plan for this session

- Multiple layers ~ Deep: MLP
- Learning rate
- More regularization
- Hyperparameters
- Backpropagation and gradients


Computing gradients over graph

SGD requires gradient of J ∇J

- For each example x_i $\frac{\delta J_{x_i}}{\delta w}$
- For each parameter $w \in W$

Backpropagation!

- Forward pass to compute J
- Backward pass to compute ∇J




Computing gradients over graph

In the backward pass, we apply the chain rule to compute gradient

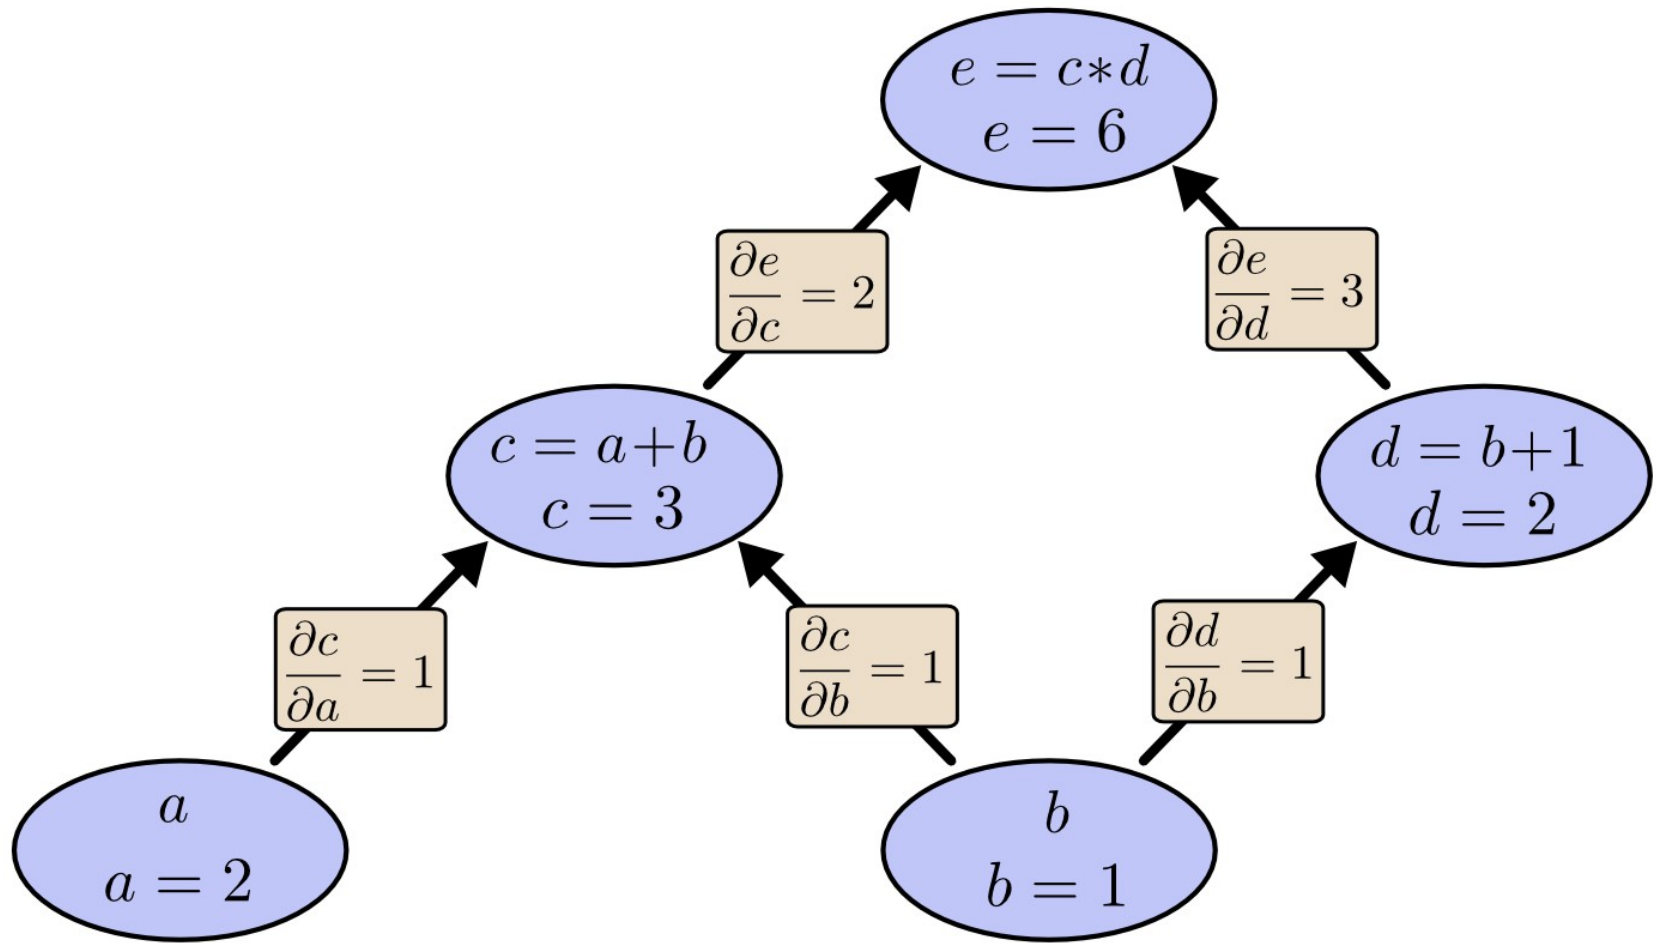
- $(f \circ g)' = (f' \circ g) g'$
- $y = f(u) \quad u = g(x)$

$$\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$$

Very efficient!

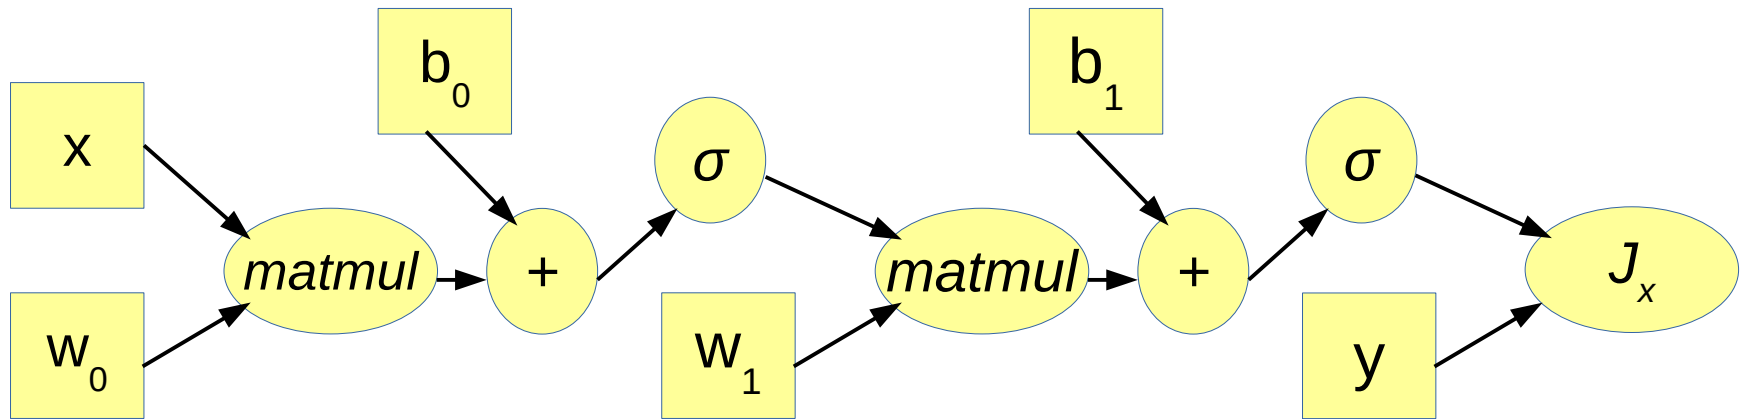


Computing gradients over graph



Source: colah.github.io/posts/2015-08-Backprop/

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

$$h_1 = \sigma(h_0 w_1 + b_1) \quad \underline{h_1}$$

$$h_0 = \sigma(x w_0 + b_0) \quad \underline{h_0}$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph

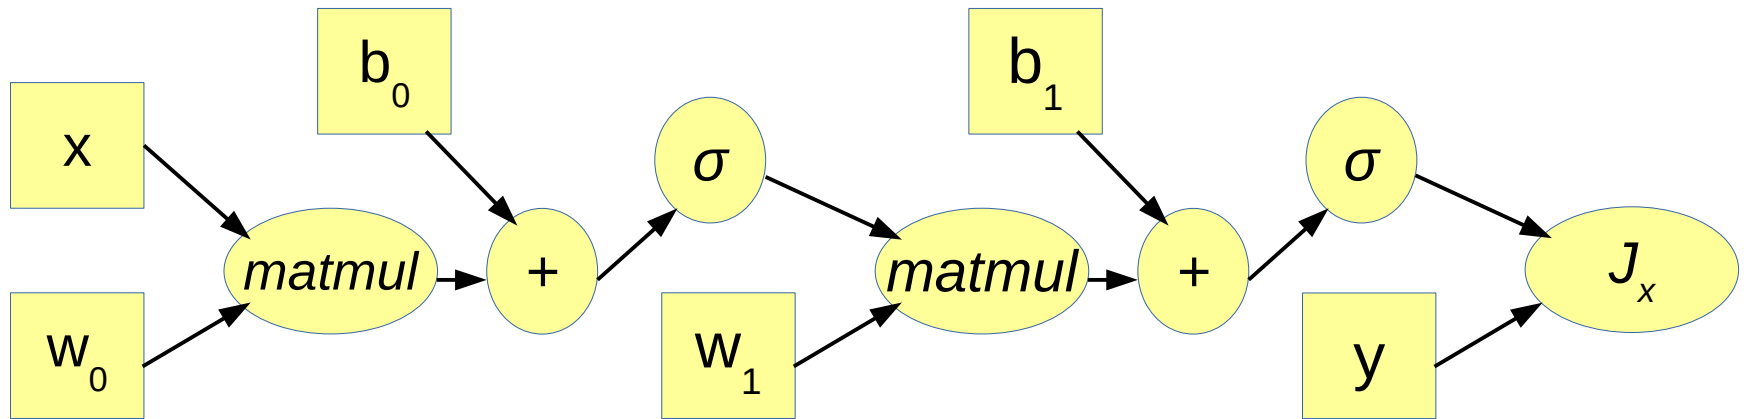
$$\begin{aligned} J_x &= -\log P(y=c|x) && \text{Cross-entropy loss} \\ &= -\log(\exp(\underline{h_1}[c]) / \sum_{c' \in C} \exp(\underline{h_1}[c'])) && \text{multiclass} \\ &= -\log(\exp(\underline{h_1}[c]) / \exp(\underline{h_1}[1]) + \exp(\underline{h_1}[0])) && y \text{ in } \{0,1\} \\ &= -\log(\exp(\underline{h_1}[c]) / Z) \\ &= -y \log(\exp(\underline{h_1}[1]) / Z) - (1-y) \log(\exp(\underline{h_1}[0]) / Z) \end{aligned}$$

$$J_x = -y \log \sigma(\underline{h_1}) - (1-y) \log(1 - \sigma(\underline{h_1})) \quad \begin{array}{l} \text{Cross-entropy} \\ \text{binary} \\ y \text{ in } \{0,1\} \end{array}$$

Chris Yeh: <https://chrisyeh96.github.io/2018/06/11/logistic-regression>



Computing gradients over graph



$$J_x = -\log P(y=c|x)$$
$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

$$h_1 = \sigma(h_0 w_1 + b_1)$$

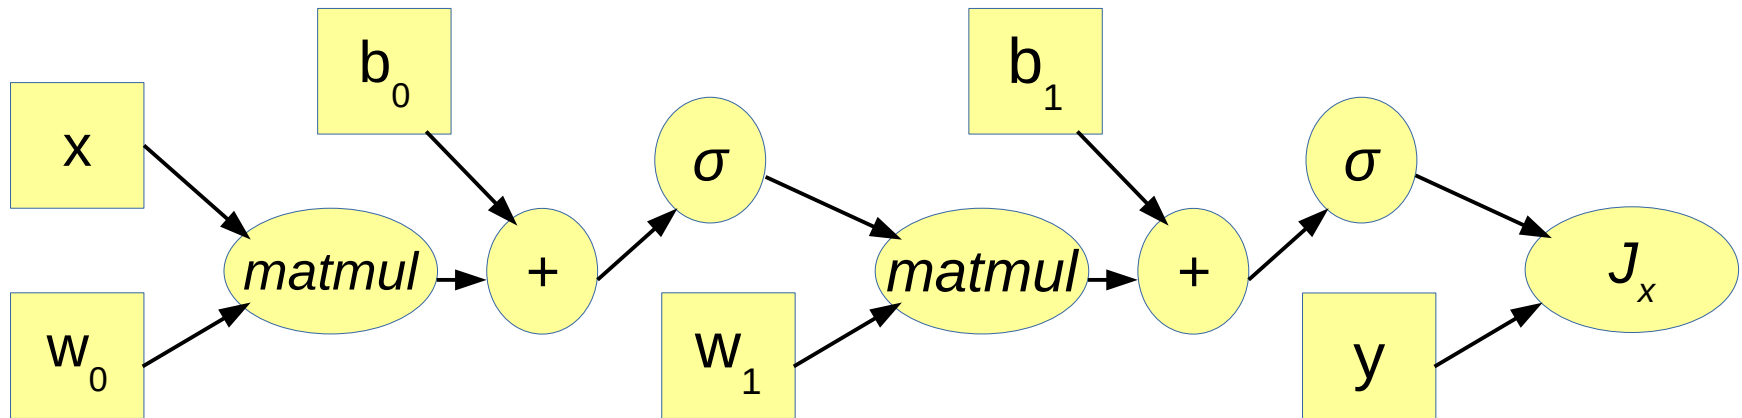
$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)



Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

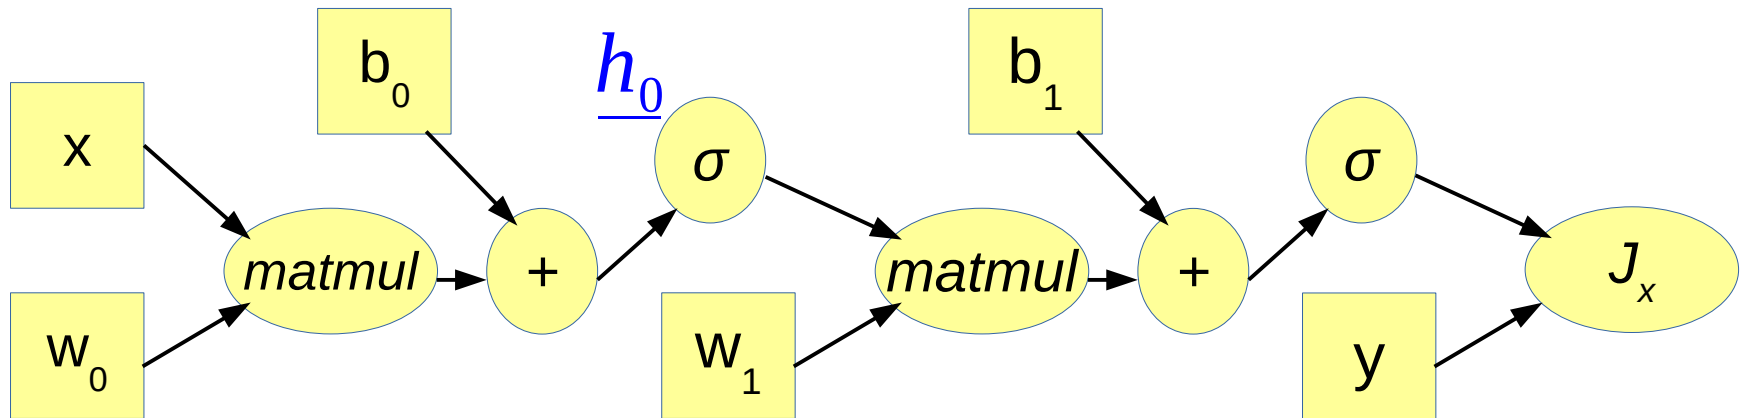
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

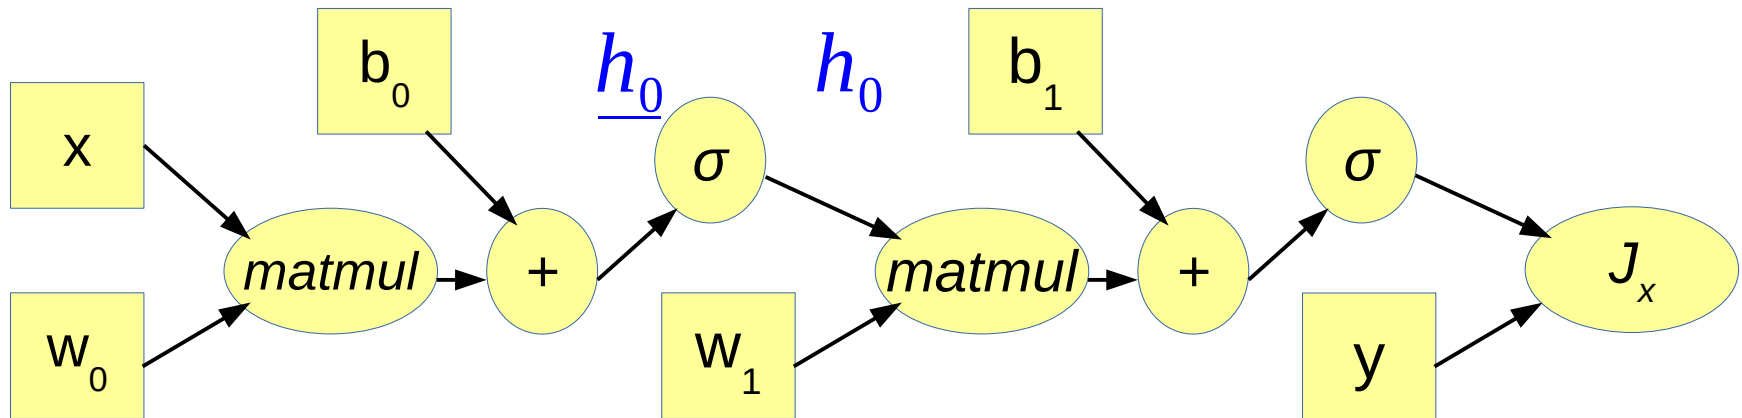
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

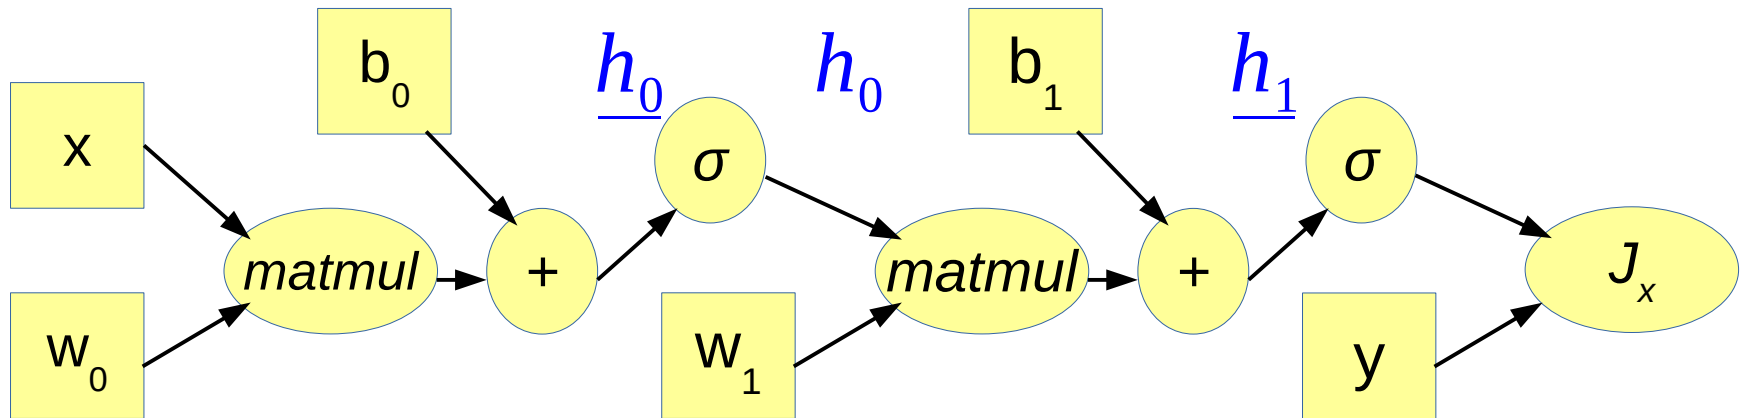
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

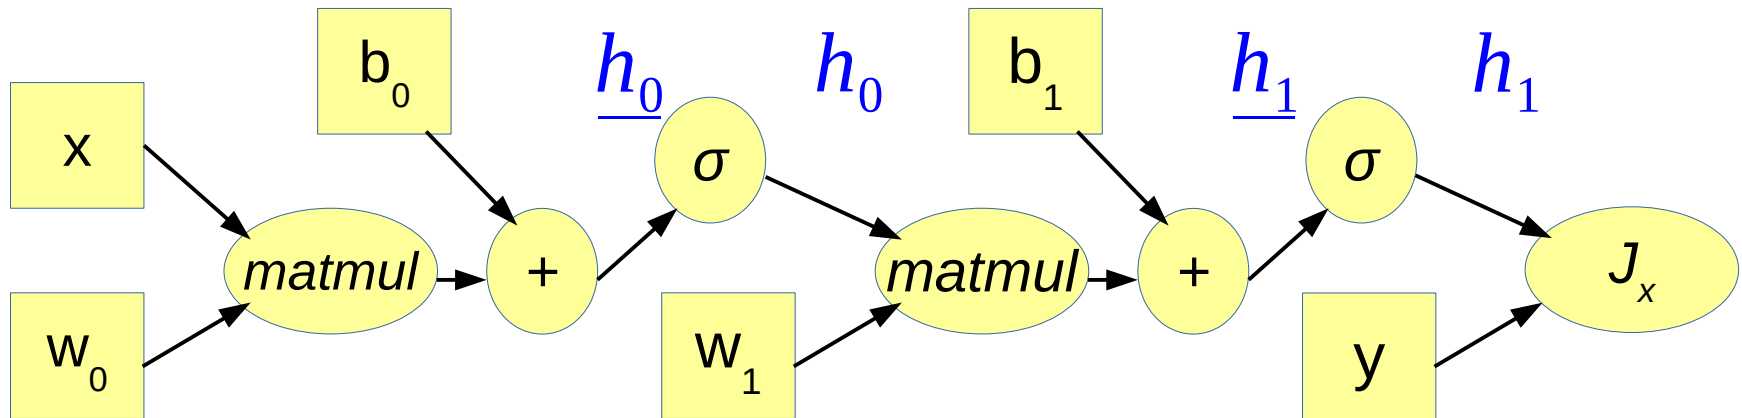
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

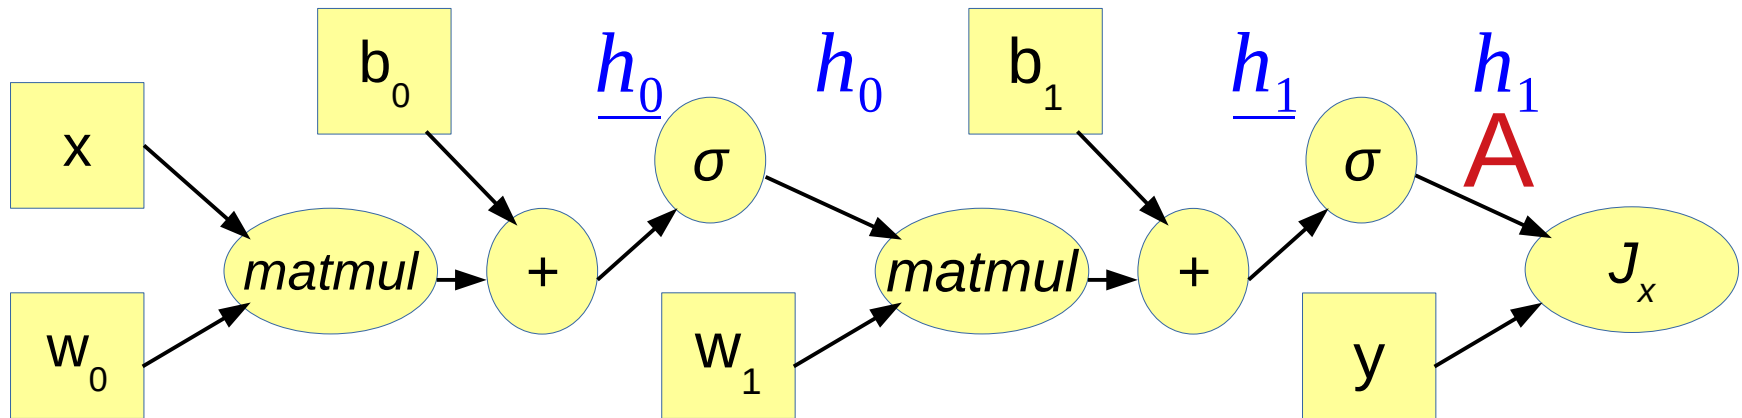
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

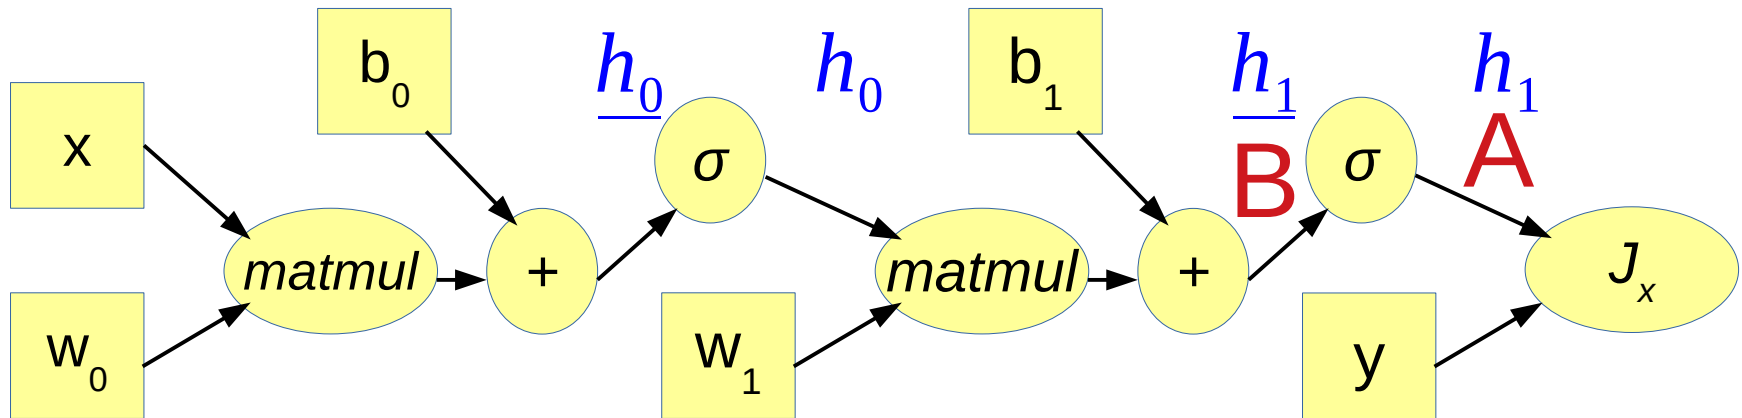
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

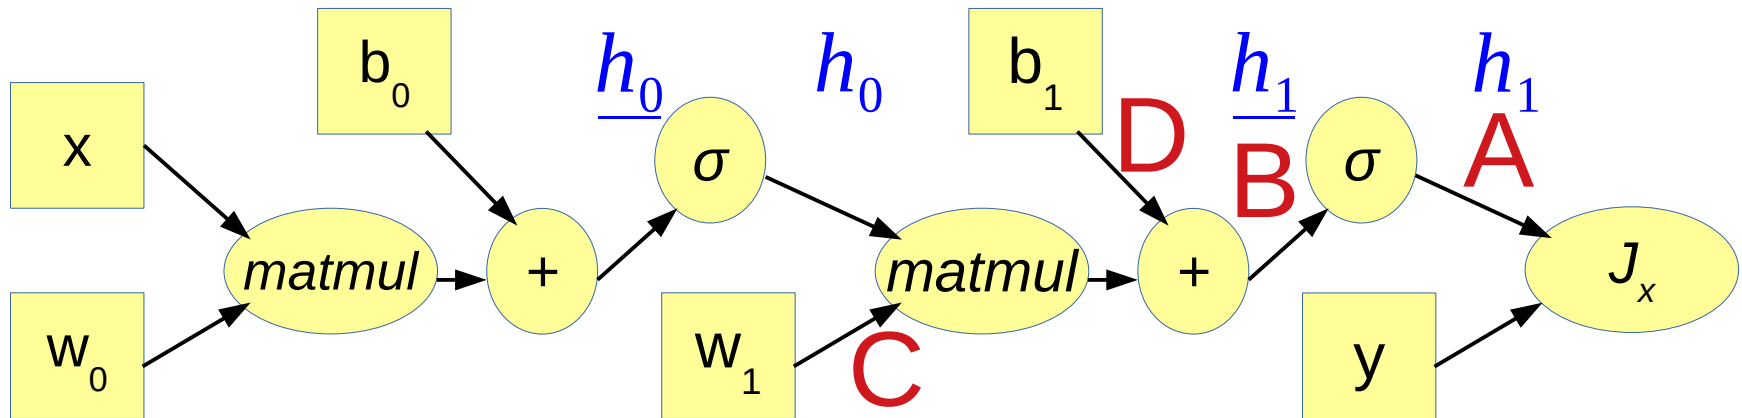
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

$$h_1 = \sigma(h_0 w_1 + b_1)$$

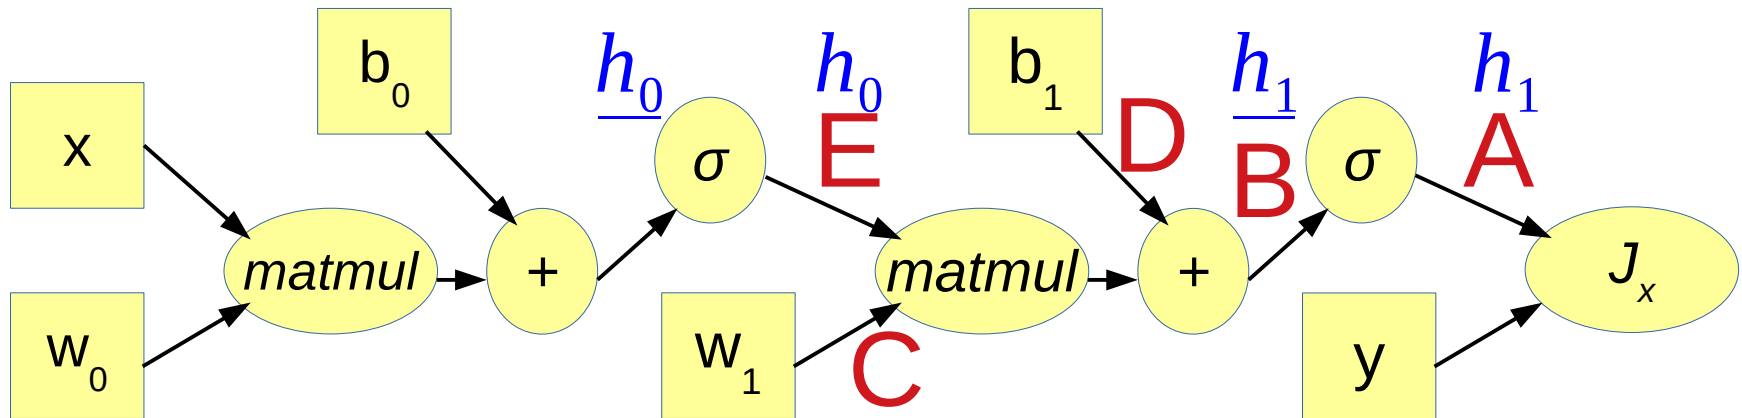
$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)



Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

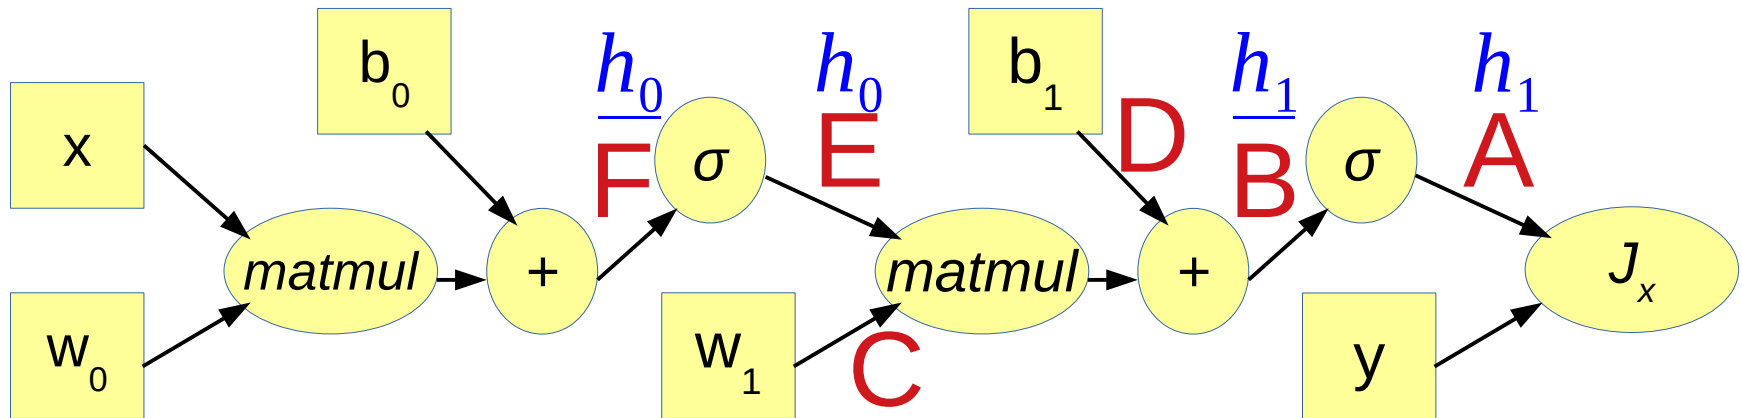
$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$

$$h_1 = \sigma(h_0 w_1 + b_1)$$

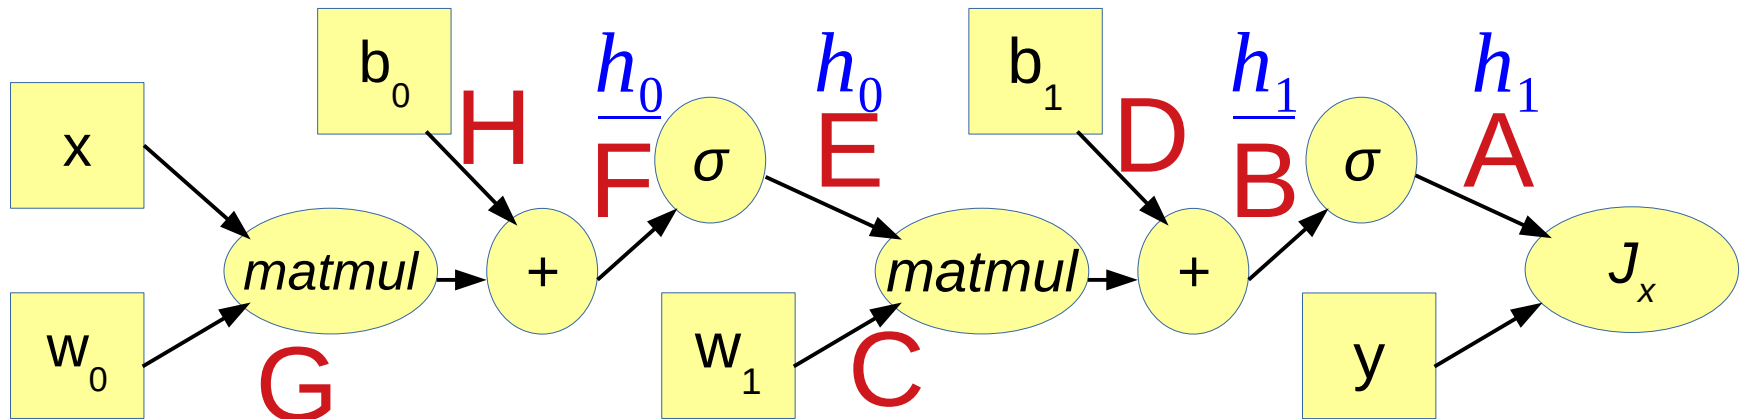
$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)



Computing gradients over graph



$$J_x = -\log P(y=c|x)$$

$$= -y \log h_1 - (1-y) \log (1-h_1)$$

Cross-entropy loss
for binary case
 y in $\{0,1\}$


$$h_1 = \sigma(h_0 w_1 + b_1)$$

$$h_0 = \sigma(x w_0 + b_0)$$

Simplification: one feature,
scalar variables

Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph



Kyunghyun Cho: https://github.com/nyu-dl/NLP_DL_Lecture_Note/raw/master/lecture_note.pdf (p. 19)

Computing gradients over graph


All operations in the graph need to be *differentiable*!

- matrix multiplication/addition
- log, exp
- sigmoid, tanh, ReLU (really!)

Many functions are not

- IF, accuracy, BLEU, etc.

=> Neural Turing Machine (differentiable)



THANKS!

Acknowledgements:

- Overall slides: Sam Bowman and Kyunghyun cho (NYU), Chris Manning and Richard Socher (Stanford)
- All source url's listed in the slides.

